

Week 03 Knowledge Check — TPM Academy

Week 3 Knowledge Check — Requirements Engineering & Mini-Capstone

A short retention check covering the week's core ideas: drafting technical PRDs, granular acceptance criteria, non-functional requirements, PRD assembly, and structured peer review. Answer each question, then check yourself against the key at the end. Aim for concepts, not trivia — every answer maps back to something we used in a lab.

Format: 12 questions (multiple choice + true/false). ~15 minutes. Individual or triad.

Questions

- 1. (MC)** What is the “TPM contract” — the test that distinguishes a Technical PRD from a generic PM PRD?
 - A. A PRD the executive sponsor approves without edits
 - B. A PRD an engineering lead would accept as scoping input *without a clarifying call*
 - C. A PRD that fits on a single page
 - D. A PRD marketing can turn into launch copy
- 2. (T/F)** Week 3 encourages using AI to draft the PRD faster, so triads can spend more of their time on review.
- 3. (MC)** Which of the following *belongs* in the §5 Solution Sketch, rather than being cut as over-specification?
 - A. Database schema decisions
 - B. Specific API contracts
 - C. The user-visible flow (4–8 steps)
 - D. Class names and library picks
- 4. (MC)** An acceptance criterion reads: “*Then dispatchers reconcile faster.*” Which of the five AC failure modes is this?

- A. Vague
- B. Untestable
- C. Restating the goal
- D. AND-soup

5. (T/F) Weird-path AC coverage — network drops, race conditions, timeouts, boundary values — is described as the TPM differentiator that PMs and pure engineers each tend to miss.

6. (MC) In the Requirement / Defense / Verification NFR form, which field is singled out as the one that separates a real TPM NFR from boilerplate?

- A. Requirement
- B. Defense
- C. Verification
- D. Category

7. (MC) How should a performance NFR express its latency target?

- A. The average (mean) response time
- B. A p95 target on a representative device and connection
- C. A best-case time under ideal conditions
- D. “Fast enough that users don’t complain”

8. (T/F) A mature NFR section includes a “Known Trade-offs” subsection that names the tensions — and their resolution or open status — rather than pretending everything is achievable at once.

9. (MC) In §10 Dependencies, what makes a dependency “owned”?

- A. It’s assigned to the responsible team
- B. It’s assigned to a named person
- C. Its status is marked “Confirmed”
- D. It has a “by when” date

10. (T/F) In §9, listing “no risks” is acceptable as long as the feature is genuinely low-risk.

11. (MC) On the Friday review rubric, which dimension carries the highest weight?

- A. Problem clarity (20%)
- B. AC testability (25%)
- C. NFR completeness (20%)
- D. Strategy linkage (15%)

12. (MC) Which reviewer rule is enforced during AC cross-review and again on Friday?

- A. Always give three specific strengths before any criticism
 - B. If you say “this could be more specific,” you also write the specific version
 - C. Never score below 2.0 without facilitator sign-off
 - D. Reviewers must not coordinate with each other
-

Answer Key

- 1. B** — The TPM contract: a PRD an engineering lead would accept as scoping input *without a clarifying call*. This is the single best evaluation criterion, and it's the bar used again on Friday.
- 2. False** — Week 3 is deliberately non-AI. AI fills vagueness with confident, generic prose — and vagueness is exactly the bug the week is built to eliminate.
- 3. C** — The user-visible flow (4–8 steps). §5 gives an engineer the *shape*; schema, API contracts, and class names are engineering judgment and are explicitly excluded.
- 4. C** — Restating the goal. It pins to a metric, not an in-system event. Fix by anchoring to an observable system state (“could you write a test without running an A/B?”).
- 5. True** — Weird-path coverage is where TPMs earn their seat. PMs who don't think technically miss it; engineers who don't think about users miss it differently.
- 6. B** — Defense. “p95 < 400ms because dispatchers tap this 40 times per shift” beats “should be fast.” The Defense field is the most-skipped part of NFRs in the wild.
- 7. B** — A p95 target on a representative device/connection. Averages lie: a bimodal 200ms/5s split averages to 2.6s but describes no real user's experience.
- 8. True** — The “Known Trade-offs” subsection is the senior-author signal. “We accept this” is strong; “we will change it if X happens” is stronger.
- 9. B** — A named person. “The auth team” is unowned; “Lin from auth” is owned. “TBD” as owner is the dependency that ships late.
- 10. False** — “No risks” is a hard fail — it's the tell of a PRD written to avoid scrutiny. Each risk gets likelihood, impact, and a mitigation.
- 11. B** — AC testability, at 25%, is the highest-weighted dimension. Scores ≥ 3.0 ship as-is; 2.0–2.9 ship with named gaps; <2.0 triggers facilitator intervention.
- 12. B** — “If you say it could be more specific, you write the specific version.” The rule eliminates lazy review feedback. (Three strengths and no reviewer coordination are also real Friday rules — but *this* is “the reviewer rule.”)