

AI Fundamentals & Prompting for TPMs

Week 1 • Day 1 • Customer-Centric Foundations

Technical Product Manager Academy



About Me

John Kidd



About Me



John Kidd

- 20+ years building commercial software

About Me



John Kidd

- 20+ years building commercial software
- 20+ years teaching technology and related subjects

About Me



John Kidd

- 20+ years building commercial software
- 20+ years teaching technology and related subjects
- First tech I taught: Lotus Notes

About Me



John Kidd

- 20+ years building commercial software
- 20+ years teaching technology and related subjects
- First tech I taught: Lotus Notes
- Yes, I love Goldens!



Prerequisites

This academy assumes you:

Please note: if any of these is your first exposure, flag it now so we can pair exercises and pace appropriately.



Prerequisites

This academy assumes you:

- Completed the pre-work — PM Philosophies, Agile Methodologies, and Prompt Engineering

Please note: if any of these is your first exposure, flag it now so we can pair exercises and pace appropriately.



Prerequisites

This academy assumes you:

- Completed the pre-work — PM Philosophies, Agile Methodologies, and Prompt Engineering
- Are comfortable with browser-based tools and a markdown editor

Please note: if any of these is your first exposure, flag it now so we can pair exercises and pace appropriately.



Prerequisites

This academy assumes you:

- Completed the pre-work — PM Philosophies, Agile Methodologies, and Prompt Engineering
- Are comfortable with browser-based tools and a markdown editor
- Have working access to a generative AI assistant (ChatGPT, Claude, or equivalent)

Please note: if any of these is your first exposure, flag it now so we can pair exercises and pace appropriately.



Where This Fits

Adjacent skills you'll touch

- Prompt engineering & AI-augmented research
- System design & architecture trade-offs
- Stakeholder mapping & negotiation
- Agile delivery and ADO mastery

Where to go next

- Capstone artifact carries into your real product work
- Ongoing TPM mentorship inside Expeditors
- Skills/SkillIQ post-academy assessments



My Pledge to You

I will...

If you have an accessibility need, please let me know — in chat, by email, or directly — and I'll adapt.



My Pledge to You

I will...

- Make this academy interactive — small-group breakouts every block

If you have an accessibility need, please let me know — in chat, by email, or directly — and I'll adapt.



My Pledge to You

I will...

- Make this academy interactive — small-group breakouts every block
- Ask you questions — and welcome yours

If you have an accessibility need, please let me know — in chat, by email, or directly — and I'll adapt.

My Pledge to You

I will...

- Make this academy interactive — small-group breakouts every block
- Ask you questions — and welcome yours
- Make sure every voice in the cohort gets heard

If you have an accessibility need, please let me know — in chat, by email, or directly — and I'll adapt.



My Pledge to You

I will...

- Make this academy interactive — small-group breakouts every block
- Ask you questions — and welcome yours
- Make sure every voice in the cohort gets heard
- Create an inclusive learning environment

If you have an accessibility need, please let me know — in chat, by email, or directly — and I'll adapt.



My Pledge to You

I will...

- Make this academy interactive — small-group breakouts every block
- Ask you questions — and welcome yours
- Make sure every voice in the cohort gets heard
- Create an inclusive learning environment
- Use an on-screen timer for breaks and lab timeboxes

If you have an accessibility need, please let me know — in chat, by email, or directly — and I'll adapt.



How We'll Work Together



How We'll Work Together

- All materials live in the academy repo — decks under week-XX/presentations/, lab briefs under week-XX/labs/



How We'll Work Together

- All materials live in the academy repo — decks under week-XX/presentations/, lab briefs under week-XX/labs/
- We'll use a generative AI assistant and a markdown editor — please install ahead per setup/TECH-SETUP.md



How We'll Work Together

- All materials live in the academy repo — decks under week-XX/presentations/, lab briefs under week-XX/labs/
- We'll use a generative AI assistant and a markdown editor — please install ahead per setup/TECH-SETUP.md
- Follow along in the day's lab brief



How We'll Work Together

- All materials live in the academy repo — decks under week-XX/presentations/, lab briefs under week-XX/labs/
- We'll use a generative AI assistant and a markdown editor — please install ahead per setup/TECH-SETUP.md
- Follow along in the day's lab brief
- Activities and labs are embedded in each day — come prepared to work in triads



How We'll Work Together

- All materials live in the academy repo — decks under week-XX/presentations/, lab briefs under week-XX/labs/
- We'll use a generative AI assistant and a markdown editor — please install ahead per setup/TECH-SETUP.md
- Follow along in the day's lab brief
- Activities and labs are embedded in each day — come prepared to work in triads
- Ask questions any time — chat, raised hand, or unmute

Cohort Introductions

In ~30 seconds each, please share:

Cohort Introductions

In ~30 seconds each, please share:

- Your name and pronouns

Cohort Introductions

In ~30 seconds each, please share:

- Your name and pronouns
- Current role and team

Cohort Introductions

In ~30 seconds each, please share:

- Your name and pronouns
- Current role and team
- Where you're based

Cohort Introductions

In ~30 seconds each, please share:

- Your name and pronouns
- Current role and team
- Where you're based
- A product you ship to today — internal or external

Cohort Introductions

In ~30 seconds each, please share:

- Your name and pronouns
- Current role and team
- Where you're based
- A product you ship to today — internal or external
- A fun fact — one sentence



Recording & Notetaker Policy

Notetaker apps (Read AI, Otter, Fireflies, etc.) are not permitted in this session.



Recording & Notetaker Policy

Notetaker apps (Read AI, Otter, Fireflies, etc.) are not permitted in this session.

- The instructor will decline notetaker requests — no exceptions



Recording & Notetaker Policy

Notetaker apps (Read AI, Otter, Fireflies, etc.) are not permitted in this session.

- The instructor will decline notetaker requests — no exceptions
- These tools capture more than shared text; they raise privacy and IP concerns



Recording & Notetaker Policy

Notetaker apps (Read AI, Otter, Fireflies, etc.) are not permitted in this session.

- The instructor will decline notetaker requests — no exceptions
- These tools capture more than shared text; they raise privacy and IP concerns
- For follow-up questions, reach out directly to the instructor



Recording & Notetaker Policy

Notetaker apps (Read AI, Otter, Fireflies, etc.) are not permitted in this session.

- The instructor will decline notetaker requests — no exceptions
- These tools capture more than shared text; they raise privacy and IP concerns
- For follow-up questions, reach out directly to the instructor

Thanks for your cooperation.



Let's Get Started!

Day 1 — AI Fundamentals & Prompting

Day 1 Objectives

Day 1 Objectives

- Describe what a large language model (LLM) can and *cannot* reliably do for a TPM

Day 1 Objectives

- Describe what a large language model (LLM) can and *cannot* reliably do for a TPM
- Apply four prompt patterns: **Role, Context, Constraints, Format**

Day 1 Objectives

- Describe what a large language model (LLM) can and *cannot* reliably do for a TPM
- Apply four prompt patterns: **Role, Context, Constraints, Format**
- Use AI to accelerate three recurring PM tasks: research, summarization, critique

Day 1 Objectives

- Describe what a large language model (LLM) can and *cannot* reliably do for a TPM
- Apply four prompt patterns: **Role, Context, Constraints, Format**
- Use AI to accelerate three recurring PM tasks: research, summarization, critique
- Detect and recover from hallucination, grounding failures, and false confidence

Day 1 Objectives

- Describe what a large language model (LLM) can and *cannot* reliably do for a TPM
- Apply four prompt patterns: **Role, Context, Constraints, Format**
- Use AI to accelerate three recurring PM tasks: research, summarization, critique
- Detect and recover from hallucination, grounding failures, and false confidence
- Build a personal **Prompt Pattern Library** with your cohort triad



Today's Shape

Clock	Block	Outcome
09:00 – 09:15	Opening	Meet your triad, ground the week
09:15 – 10:30	Block 1 + Activity 1	What LLMs actually are
10:45 – 12:00	Block 2 + Activity 2	Four prompt patterns
13:00 – 14:15	Block 3 + Activity 3	Three TPM uses of AI
14:30 – 16:00	Block 4 + Activity 4 + Readouts	Verification & trust

Block 1 — What an LLM Actually Is

Enough mental model to stop being surprised by it

The Only Mental Model You Need

An LLM is a **very large statistical model of language** that predicts the next token given everything before it. It is not a database. It is not a reasoner. It does not "know" things.

It is good at

- Rephrasing, summarizing, re-formatting
- Pattern-completing from strong context
- Drafting first passes quickly
- Stress-testing your own writing

It is bad at

- Giving exact numbers or citations from memory
- Saying "I don't know"
- Reasoning about novel constraints without being shown them

Three Terms You Will Say Out Loud

Token

The unit the model reads and writes.
Roughly ~4 characters of English. "Product" $\approx$ 1 token, "Dispatch" $\approx$ 1–2.

Context window

How many tokens the model can see in one call.
Everything outside this window is invisible to it.

Temperature

A knob from deterministic (0) to creative (1+). PMs usually want low-to-mid for drafts, low for extraction.

PM implication: If the model can't see it in the prompt, it's guessing. Always ask yourself: *"What does this model know about my product right now? Nothing."*

Activity 1 — "Same Question, Three Models"

Triads • 35 minutes

Each triad picks one open PM question (provided list of 10). Each member asks the *same* question to a different tool (ChatGPT, Claude, Gemini/Copilot).

- Compare: Where do the answers agree? Diverge? Hallucinate?
- Score each answer on **usefulness**, **verifiability**, **risk if wrong**
- Prepare a 60-second readout: "We would trust X for Y, and never trust it for Z"



10 min prompt · 15 min compare · 10 min write readout

Block 2 — Four Prompt Patterns

The smallest scaffold that meaningfully improves output

The RCCF Scaffold

Every time you ask an AI for non-trivial PM help, include four things:
Role, Context, Constraints, Format.

Slot	Question it answers	PM example
Role	Who is talking?	"Act as a senior technical product manager reviewing a PRD"
Context	What does the model need to see?	User persona, the problem, prior attempts
Constraints	What must be true of the answer?	"Flag any claim that lacks evidence"
Format	What shape should it take?	"Markdown table with columns X, Y, Z"



Before & After

✗ Weak prompt

Write me a user persona for a dispatch app.

Output is generic, confident, ungrounded, un-actionable.

✓ Scaffolded prompt

Role: Senior TPM for a field-service SaaS.
Context: Target user = dispatcher at an HVAC company, 50–200 techs, 10+ yrs industry, low tolerance for screens. Pain: re-routing mid-day when a tech calls out sick.
Constraints: Persona must be grounded in the behaviors above; flag any field that is an assumption I'd need to validate.
Format: Markdown persona sheet with a final "Assumptions to validate" section.



Prompting is a Dialogue, Not a Command

Caution: The model will happily invent critiques too. Your judgment is what converts dialogue into value.



Prompting is a Dialogue, Not a Command

- **First pass:** RCCF scaffold

Caution: The model will happily invent critiques too. Your judgment is what converts dialogue into value.



Prompting is a Dialogue, Not a Command

- **First pass:** RCCF scaffold
- **Second pass:** "Critique the above as if you were a skeptical engineer"

Caution: The model will happily invent critiques too. Your judgment is what converts dialogue into value.



Prompting is a Dialogue, Not a Command

- **First pass:** RCCF scaffold
- **Second pass:** "Critique the above as if you were a skeptical engineer"
- **Third pass:** "Now rewrite it, addressing those critiques"

Caution: The model will happily invent critiques too. Your judgment is what converts dialogue into value.



Prompting is a Dialogue, Not a Command

- **First pass:** RCCF scaffold
- **Second pass:** "Critique the above as if you were a skeptical engineer"
- **Third pass:** "Now rewrite it, addressing those critiques"
- **Always:** "What in this answer are you least confident about?"

Caution: The model will happily invent critiques too. Your judgment is what converts dialogue into value.

Activity 2 — Rewrite This Prompt

Triads • 40 minutes

Each triad receives three weak prompts taken from the FieldPulse backlog. Rewrite each one using RCCF, then run the original and the rewrite side-by-side.

- Capture: What changed in the output? What did *not* change?
- Add one prompt to your triad's **Prompt Pattern Library**
- Star the prompt your triad would actually reuse on Monday morning

 15 min rewrite · 15 min run · 10 min library entry

Block 3 — Three Recurring TPM Uses

Where AI actually earns its seat at the table



Use 1 — Accelerated Research

The job: Quickly build context about an unfamiliar domain, user group, or competitor landscape — enough to ask smart questions, not to ship decisions.

✓ Good fit

- "Explain HVAC dispatcher workflows to me in the language they'd use"
- "What are five things I might be naive about here?"
- "Give me five search queries that will find primary-source material"

✗ Bad fit

- "What is the market size of dispatch software?" (will hallucinate)
- "Who is the leader?" (outdated and confident)
- "Cite a study that proves..." (fabricated citations)



Use 2 — Summarization with Evidence Tags

When you feed a model a transcript, survey, or support-ticket batch, ask it to **tag each claim with where in the source it came from**. Unsourced claims are the ones to interrogate.

```
Role: Research analyst.  
Context: Below are 14 support tickets from FieldPulse dispatchers.  
Task: Extract the top 5 recurring frustrations.  
Constraints:  
- Each frustration must cite 2+ ticket numbers as evidence  
- Flag any frustration you inferred without direct evidence  
Format: Markdown list; each entry: Frustration / Evidence tickets / Confidence (H/M/L).
```

Use 3 — Self-Critique Loop

Before sharing a PRD, persona, or problem statement, run a structured critique pass:

1. "Read the following as a skeptical engineer. What three questions will you ask me first?"
2. "Read the following as the stakeholder who most often blocks launches here. What objection worries me most?"
3. "Read the following as the customer in the persona. Where does it not ring true?"

Why it works: You get three perspectives in five minutes without socializing a half-baked draft with real people.

Activity 3 — The Three Hats

Triads • 40 minutes

Each triad takes the FieldPulse one-paragraph problem brief (provided). Using AI, run it through three critique hats: **Engineer**, **Blocker-stakeholder**, **Target customer**.

- Capture the single most uncomfortable question each hat raised
- Decide which one your triad cannot answer yet — that becomes your Day 3 research question



10 min each hat • 10 min synthesize

Block 4 — Verification & Trust

How grown-up TPMs use AI without being fooled by it



The Four Failure Modes

Failure	What it looks like	Mitigation
Hallucination	Confident fabricated facts, citations, APIs	Treat outputs as leads, not sources. Verify numerics & names manually.
Context starvation	Plausible but generic answer	Add Role + Context explicitly. Paste in the source material.
Sycophancy	Agrees with whatever premise you give	Ask for counter-arguments and rank them by strength.
Anchoring	Your framing shapes its output	Ask for the opposite position as a second turn.

1 2
3 4

Evidence Tiers (we'll use this all week)

When you write something down — persona traits, pain points, problem statements — tag each claim with its evidence tier so reviewers can calibrate trust.

Tier	Source	Trust
Observed	You watched it happen, or recorded data shows it	High
Reported	A user said it in an interview or ticket	Medium
Inferred	You deduced it from related evidence	Low
AI-generated	A model produced it — not yet validated	Lowest; must be upgraded before use

Activity 4 — Prompt Pattern Library (Publish)

Triads • 45 minutes • Readouts

Each triad finalizes a shared **Prompt Pattern Library** for use the rest of the academy. Minimum four entries:

- One **Research** prompt
- One **Summarization-with-evidence** prompt
- One **Critique-hat** prompt
- One **"Where could you be wrong?"** prompt

 25 min draft • 10 min test • 10 min triad readout



Day 1 Wrap

What we built

- A mental model of what an LLM is (and isn't)
- The RCCF prompt scaffold
- Three reusable TPM uses: research, summarization, critique
- A Prompt Pattern Library per triad

What opens tomorrow

- **Working Backwards** — we write the press release *before* the product
- Your triad will use today's prompts to draft and critique your first PR/FAQ

Bring to Day 2: Your triad's Prompt Pattern Library + one domain you'd like to draft a Working-Backwards PR for.

Questions & Reflection

Before you leave: what's the one prompt you'll use tomorrow morning?