



AI Fundamentals & Prompting for TPMs

Week 1 • Day 1 • Customer-Centric Foundations

Technical Product Manager Academy





Let's Get Started!

Day 1 — AI Fundamentals & Prompting

Day 1 Objectives

- Describe what a large language model (LLM) can and *cannot* reliably do for a TPM
- Apply four prompt patterns: **Role, Context, Constraints, Format**
- Use AI to accelerate three recurring PM tasks: research, summarization, critique
- Detect and recover from hallucination, grounding failures, and false confidence
- Build a personal **Prompt Pattern Library** with your cohort quad



Today's Shape

Clock	Block	Outcome
09:00 – 09:15	Opening	Meet your quad, ground the week
09:15 – 10:30	Block 1 + Activity 1	What LLMs actually are
10:45 – 12:00	Block 2 + Activity 2	Four prompt patterns
13:00 – 14:15	Block 3 + Activity 3	Three TPM uses of AI
14:30 – 16:00	Block 4 + Activity 4 + Readouts	Verification & trust

Block 1 — What an LLM Actually Is

Enough mental model to stop being surprised by it

The Only Mental Model You Need

An LLM is a **very large statistical model of language** that predicts the next token given everything before it. It is not a database. It is not a reasoner. It does not "know" things.

It is good at

- Rephrasing, summarizing, re-formatting
- Pattern-completing from strong context
- Drafting first passes quickly
- Stress-testing your own writing

It is bad at

- Giving exact numbers or citations from memory
- Saying "I don't know"
- Reasoning about novel constraints without being shown them
- Telling you when it is wrong



Three Terms You Will Say Out Loud

Token

The unit the model reads and writes.

Roughly ~4 characters of English. "Product" $\approx$ 1 token, "Dispatch" $\approx$ 1–2.

Context window

How many tokens the model can see in one call. Everything outside this window is invisible to it.

Temperature

A knob from deterministic (0) to creative (1+). PMs usually want low-to-mid for drafts, low for extraction.

PM implication: If the model can't see it in the prompt, it's guessing. Always ask yourself: *"What does this model know about my product right now? Nothing."*

Activity 1 — "Same Question, Three Models"

Quads • 45 minutes

Each quad picks one open PM question (provided list of 10). Each member asks the *same* question to a different tool (ChatGPT, Claude, Gemini/Copilot).

- Compare: Where do the answers agree? Diverge? Hallucinate?
- Score each answer on **usefulness**, **verifiability**, **risk if wrong**
- Prepare a 60-second readout: "We would trust X for Y, and never trust it for Z"



10 min prompt • 25 min compare • 10 min write readout

Block 2 — Four Prompt Patterns

The smallest scaffold that meaningfully improves output



The RCCF Scaffold

Every time you ask an AI for non-trivial PM help, include four things: **Role, Context, Constraints, Format**.

Slot	Question it answers	PM example
Role	Who is talking?	"Act as a senior technical product manager reviewing a PRD"
Context	What does the model need to see?	User persona, the problem, prior attempts
Constraints	What must be true of the answer?	"Flag any claim that lacks evidence"
Format	What shape should it take?	"Markdown table with columns X, Y, Z"



Before & After

✗ Weak prompt

Write me a user persona for a dispatch app.

Output is generic, confident, ungrounded, un-actionable.



✓ Scaffolded prompt

Role: Senior TPM for a field-service SaaS.
Context: Target user = dispatcher at an HVAC company, 50-200 techs, 10+ yrs industry, low tolerance for screens. Pain: re-routing mid-day when a tech calls out sick.
Constraints: Persona must be grounded in the behaviors above; flag any field that is an assumption I'd need to validate.
Format: Markdown persona sheet with a final "Assumptions to validate" section.



Prompting is a Dialogue, Not a Command

- **First pass:** RCCF scaffold
- **Second pass:** "Critique the above as if you were a skeptical engineer"
- **Third pass:** "Now rewrite it, addressing those critiques"
- **Always:** "What in this answer are you least confident about?"

Caution: The model will happily invent critiques too. Your judgment is what converts dialogue into value.

Activity 2 — Rewrite This Prompt

Quads • 50 minutes

Each quad receives three weak prompts taken from the FieldPulse backlog. Rewrite each one using RCCF, then run the original and the rewrite side-by-side.

- Capture: What changed in the output? What did *not* change?
- Add one prompt to your quad's **Prompt Pattern Library**
- Star the prompt your quad would actually reuse on Monday morning

 25 min rewrite • 15 min run • 10 min library entry

Block 3 — Three Recurring TPM Uses

Where AI actually earns its seat at the table

Use 1 — Accelerated Research

The job: Quickly build context about an unfamiliar domain, user group, or competitor landscape — enough to ask smart questions, not to ship decisions.

Good fit

- "Explain HVAC dispatcher workflows to me in the language they'd use"
- "What are five things I might be naive about here?"
- "Give me five search queries that will find primary-source material"

Bad fit


- "What is the market size of dispatch software?" (will hallucinate)
- "Who is the leader?" (outdated and confident)
- "Cite a study that proves..." (fabricated citations)



Use 2 — Summarization with Evidence Tags

When you feed a model a transcript, survey, or support-ticket batch, ask it to **tag each claim with where in the source it came from**. Unsourced claims are the ones to interrogate.

```
Role: Research analyst.  
Context: Below are 14 support tickets from FieldPulse dispatchers.  
Task: Extract the top 5 recurring frustrations.  
Constraints:  
- Each frustration must cite 2+ ticket numbers as evidence  
- Flag any frustration you inferred without direct evidence  
Format: Markdown list; each entry: Frustration / Evidence tickets / Confidence (H/M/L).
```

Use 3 — Self-Critique Loop

Before sharing a product requirements document (PRD), persona, or problem statement, run a structured critique pass:

1. "Read the following as a skeptical engineer. What three questions will you ask me first?"
2. "Read the following as the stakeholder who most often blocks launches here. What objection worries me most?"
3. "Read the following as the customer in the persona. Where does it not ring true?"

Why it works: You get three perspectives in five minutes without socializing a half-baked draft with real people.

Activity 3 — The Three Hats

Quads • 50 minutes

Each quad takes the FieldPulse one-paragraph problem brief (provided). Using AI, run it through three critique hats: **Engineer**, **Blocker-stakeholder**, **Target customer**.

- Capture the single most uncomfortable question each hat raised
- Decide which one your quad cannot answer yet — that becomes your Day 3 research question



10 min each hat • 20 min synthesize

Block 4 — Verification & Trust

How grown-up TPMs use AI without being fooled by it



The Four Failure Modes

Failure	What it looks like	Mitigation
Hallucination	Confident fabricated facts, citations, APIs	Treat outputs as leads, not sources. Verify numerics & names manually.
Context starvation	Plausible but generic answer	Add Role + Context explicitly. Paste in the source material.
Sycophancy	Agrees with whatever premise you give	Ask for counter-arguments and rank them by strength.
Anchoring	Your framing shapes its output	Ask for the opposite position as a second turn.



Evidence Tiers (we'll use this all week)

When you write something down — persona traits, pain points, problem statements — tag each claim with its evidence tier so reviewers can calibrate trust.

Tier	Source	Trust
Observed	You watched it happen, or recorded data shows it	High
Reported	A user said it in an interview or ticket	Medium
Inferred	You deduced it from related evidence	Low
AI-generated	A model produced it — not yet validated	Lowest; must be upgraded before use

Activity 4 — Prompt Pattern Library (Publish)

Quads • 55 minutes • Readouts

Each quad finalizes a shared **Prompt Pattern Library** for use the rest of the academy. Minimum four entries:

- One **Research** prompt
- One **Summarization-with-evidence** prompt
- One **Critique-hat** prompt
- One **"Where could you be wrong?"** prompt

 35 min draft • 10 min test • 10 min quad readout



Day 1 Wrap

What we built

- A mental model of what an LLM is (and isn't)
- The RCCF prompt scaffold
- Three reusable TPM uses: research, summarization, critique
- A Prompt Pattern Library per quad

What opens tomorrow

- **Working Backwards** — we write the press release *before* the product
- Your quad will use today's prompts to draft and critique your first PR/FAQ

Bring to Day 2: Your quad's Prompt Pattern Library + one domain you'd like to draft a Working-Backwards PR for.

Questions & Reflection

Before you leave: what's the one prompt you'll use tomorrow morning?