




Working Backwards

Week 1 • Day 2 • Customer-Centric Foundations

Technical Product Manager Academy



Day 2 Objectives

- Explain why **Working Backwards** out-performs feature-first planning
- Structure and draft a one-page **PR/FAQ** for a product idea
- Answer the **Five Customer Questions** concretely, not aspirationally
- Apply Day 1 prompts to accelerate PR/FAQ drafting and critique
- Run a structured peer critique of another quad's PR/FAQ



Today's Shape

Clock	Block	Outcome
09:00 – 09:15	Opening	Recap Day 1, set up FieldPulse stretch case
09:15 – 10:30	Block 1 + Activity 1	Forward vs Backward thinking
10:45 – 12:00	Block 2 + Activity 2	PR/FAQ structure, headline drafting
13:00 – 14:15	Block 3 + Activity 3	The Five Customer Questions
14:30 – 16:00	Block 4 + Activity 4 + Readouts	Peer critique + revision

Block 1 — Why Work Backwards?

The feature factory problem and the fix



The Feature Factory Pattern

❌ Forward-built product

- Idea → spec → build → ship
- "Customer value" asserted, not tested
- Launch reveals whether it was needed
- Re-work paid for in production

✅ Backward-built product

- Imagine the launch → write the PR → answer the FAQ
- Customer value is *written down* first
- If the PR is weak, the product is weak
- Re-work paid for in writing

The shift: We do not use documents to plan what we'll build. We use them to *decide whether to build at all*.



Where Working Backwards Came From

Amazon formalized the **PR/FAQ** process in the early 2000s. Teams wrote a press release *before* the product existed, answered the FAQs an exec would ask, and used the document as the meeting agenda for "should we do this?"

Why it spread: It reliably kills weak ideas early, surfaces the hard questions on paper (where answers are cheap), and forces alignment before a single engineer is staffed.

Activity 1 — Forward or Backward?

Quads • 45 minutes

Each quad receives 4 short vignettes describing how a real (anonymized) product was conceived. Decide: was this forward-built or backward-built? What's the evidence?

- Two of the four are misleading — read carefully
- For each, write the *one* question that, if asked earlier, would have changed the outcome



15 min analyze • 15 min full-room debrief

Block 2 — The PR/FAQ Structure

Anatomy of the document that kills bad ideas



One-Page PR — Six Parts

#	Section	Purpose
1	Heading	Product + audience in one line a customer would read
2	Sub-heading	Concrete benefit in customer language
3	Summary paragraph	The story, in 3-4 sentences, as if launched
4	Problem paragraph	The pain before your product existed
5	Solution paragraph	How your product removes the pain — mechanism explicit
6	Customer quote	A fictional but credible customer — if you can't write it believably, the PR fails



A FieldPulse PR — Working Sketch

FieldPulse launches "End-of-Day Reconcile" — 5-minute close-outs for HVAC dispatchers

Small-business HVAC dispatchers can now close their day in five minutes instead of 45, using a guided reconcile flow that pulls truck inventory, tech timecards, and parts usage into one approval screen.

Dispatchers at 50-to-200-technician shops typically stay 45 minutes past close every night reconciling paper tickets, truck stock, and timecards. FieldPulse's new End-of-Day Reconcile pulls these into a single guided flow, flags anomalies, and submits the day in under five minutes.

"I used to leave after my kids were in bed. Last week I made it home for dinner four nights in a row." — Maria R., Lead Dispatcher, Reliable Climate Co.



The FAQ Side

Customer FAQ

- How much does it cost?
- What if my truck inventory is wrong?
- Can I use it offline?
- How does it know what parts were used?

Internal FAQ

- What's the riskiest assumption?
- What's the build estimate? On what evidence?
- What breaks if adoption is 30% of plan?
- What existing integrations block this?

Rule: If you can't write the internal FAQ, you don't understand the product yet. Go back to the PR.

Activity 2 — Headline + Sub-headline

Quads • 50 minutes

Each quad selects one product idea (from their Day 1 critique artifacts, or a provided list). Draft **three** candidate heading + sub-headings. Then use AI to critique each:

- "Which of these would a dispatcher actually click on?"
- "Which promises something we can't deliver?"
- "Which is interchangeable with a competitor's marketing?"



15 min draft • 15 min AI critique • 15 min pick + justify

Block 3 — The Five Customer Questions

The diagnostic under every PR/FAQ

? Five Questions That Force Clarity

1. **Who is the customer?** Named role, named context, not "users."
2. **What is the customer problem or opportunity?** A verb, not a category.
3. **What is the most important benefit?** One benefit. If you list three, you don't have one yet.
4. **How do you know what they want or need?** Name the evidence. Tag the tier.
5. **What does the experience look like?** Describe the moment, not the feature.

If your PR/FAQ answers all five with specifics, scoping is easy. If it doesn't, no amount of engineering rigor will save it.



Common Failure Modes

Failure	Symptom	Fix
Generic customer	"Small business owners"	Name the role, size, and context
Category problem	"Inefficient workflows"	Replace with a verb the customer uses
Benefit bouquet	"Faster, better, cheaper"	Choose one. Defend it.
Handwave evidence	"Customers told us"	Cite <i>who</i> , <i>when</i> , <i>what tier</i>
Feature description	"A dashboard with..."	Describe the human moment of use

Activity 3 — Answer the Five, Concretely

Quads • 50 minutes

Your quad has a headline/sub-headline from Activity 2. Now draft a full PR paragraph + customer quote + 3 Customer-FAQ + 3 Internal-FAQ entries.

- Use Day 1's research + summarization prompts to ground each claim
- Tag every claim with an evidence tier (**Observed** , **Reported** , **Inferred** , **AI-generated**)
- Be honest — unvalidated is fine today, but it must be *labeled*

 10 min plan • 25 min draft • 10 min self-review

Block 4 — Peer Critique & Revision

How to review without rewriting someone else's work

The Critique Protocol

Critique runs on a timed protocol so every author gets balanced feedback:

1. **2 min silent read** — reviewer highlights 1 line they love, 1 they question
2. **3 min "What I heard"** — reviewer restates the product to the author in one sentence
3. **5 min questions only** — reviewer asks, author listens and takes notes (no defending)
4. **5 min author response** — author says what they'll change, what they'll leave

Why "questions only": It forces reviewers to reveal what confused them rather than prescribing a fix. It respects the author's ownership.

Activity 4 — Cross-Quad PR/FAQ Review

Paired quads • 60 minutes • Readouts

Quads pair up. Using the critique protocol, each quad reviews the other's PR/FAQ. Then 15 minutes to revise.

- One critique from each failure-mode table cell
- Capture which evidence tier needs to be upgraded before Friday
- Each quad does a 60-second readout: *"What we changed and why"*



20 min review A • 20 min review B • 15 min revise • 5 min readout



Day 2 Wrap

What we built

- A complete one-page PR/FAQ per quad
- The Five Customer Questions as a working checklist
- A critique protocol we'll reuse every week

What opens tomorrow

- **Personas** — we take "Who is the customer?" from a sentence to a validated artifact
- Your PR/FAQ customer quote becomes persona input on Day 3

Bring to Day 3: Your PR/FAQ + the one unvalidated assumption you're most worried about.

Questions & Reflection

Which of the Five Questions was hardest for your quad?