

Strategy Brief Template

Day 4 · keeper-artifact handout. The 2–3 page brief that integrates everything from Week 1 forward. Hand-written sections carry the judgment; AI-assembled sections carry the evidence. The provenance log is mandatory — a strategy brief without it is fiction.

```
# Strategy Brief – <Problem area>
**Quad:** <names> | **NS:** <one line>

## Executive summary (5 bullets max, hand-written)
...

## Top 5 user pain themes (AI-summarized, evidence-
tagged)
For each theme:
– **Theme:** <verb-based statement>
– **Evidence:** <source citations, e.g., "Interview
A3, C2; Tickets T-04, T-11; #voice-of-customer 2025-
09-12">
– **Confidence:** H / M / L
– **AI-generated?** Y/N (if Y: which prompt + what
we did to validate)

## Competitive snapshot (AI-drafted, hand-edited)
| Competitor | Strength | Weakness | Where they beat
us |
|-----|-----|-----|-----
----|

## Strategy implications (hand-written)
– What our NS implies for these themes
– Which design principle from Day 3 most directly
addresses each theme
– One open question we still need humans for

## Provenance log
– Prompts used: link to quad's Pattern Library
entries
– Manual validation steps taken
– Where the AI was wrong (named openly)
```

What "good" looks like for the pain themes

- 5 themes (no fewer — the AI will naturally lump).
- Every theme cites at least 2 distinct sources.
- No more than 1 theme is AI-only (no observed evidence) — and it's flagged as a hypothesis to validate.
- Themes are verbs, not categories ("Cannot reconcile mid-route" beats "Reconcile UX").

The **Provenance log** is where trust is built. If it has no "where AI was wrong" entry, you didn't validate — revisit at least one theme.