

# Strategy Brief Template

**Day 4 · keeper-artifact handout.** The 2–3 page brief that integrates everything from Week 1 forward. Hand-written sections carry the judgment; AI-assembled sections carry the evidence. The provenance log is mandatory — a strategy brief without it is fiction.

```
# Strategy Brief – <Problem area>
**Triad:** <names> | **NS:** <one line>

## Executive summary (5 bullets max, hand-written)
...

## Top 5 user pain themes (AI-summarized, evidence-tagged)
For each theme:
- **Theme:** <verb-based statement>
- **Evidence:** <source citations, e.g., "Interview A3, C2; Tickets T-04, T-11;
  #voice-of-customer 2025-09-12">
- **Confidence:** H / M / L
- **AI-generated?** Y/N (if Y: which prompt + what we did to validate)

## Competitive snapshot (AI-drafted, hand-edited)
| Competitor | Strength | Weakness | Where they beat us |
|-----|-----|-----|-----|

## Strategy implications (hand-written)
- What our NS implies for these themes
- Which design principle from Day 3 most directly addresses each theme
- One open question we still need humans for

## Provenance log
- Prompts used: link to triad's Pattern Library entries
- Manual validation steps taken
- Where the AI was wrong (named openly)
```

## What "good" looks like for the pain themes

- 5 themes (no fewer — the AI will naturally lump).
- Every theme cites at least 2 distinct sources.
- No more than 1 theme is AI-only (no observed evidence) — and it's flagged as a hypothesis to validate.
- Themes are verbs, not categories ("Cannot reconcile mid-route" beats "Reconcile UX").

The **Provenance log** is where trust is built. If it has no "where AI was wrong" entry, you didn't validate — revisit at least one theme.