



# AI for Strategy & Research Summaries

Week 2 • Day 4 • Strategic Design & Metrics

TPM Academy



## Day 4 Objectives

- Use AI to **synthesize multi-source research** with traceable evidence
- Detect and correct the four **summarization failure modes**: theme inflation, citation invention, premature confidence, source bias
- Stay on the **safe side of the bright line**: structure source material, never ask AI to invent market facts
- Produce a **provenance log** that makes your AI use auditable
- Assemble an **AI-Assembled Strategy Brief** for your triad's problem



# Today's Shape

| Clock            | Block                              | Outcome                                           |
|------------------|------------------------------------|---------------------------------------------------|
| 09:00 –<br>09:15 | Opening                            | Re-ground in your Day-3 audit + design principles |
| 09:15 –<br>10:30 | Block 1 + Activity 1               | Multi-source synthesis drill                      |
| 10:45 –<br>12:00 | Block 2 + Activity 2               | The Research Pack pass                            |
| 13:00 –<br>14:15 | Block 3 + Activity 3               | Competitive snapshot                              |
| 14:30 –<br>16:00 | Block 4 + Activity 4 +<br>Readouts | Brief assembly + provenance                       |

# Block 1 — Strategy-Scale Summarization

Beyond a single document, beyond a single prompt

# What's Different at Strategy Scale

| Week 1 prompting                    | Strategy-scale prompting                     |
|-------------------------------------|----------------------------------------------|
| Single document or short brief      | 5–15 documents, often contradictory          |
| One prompt, one output              | Iterated prompts, validated output           |
| "Summarize this"                    | "Synthesize across these, tag evidence"      |
| You verify confidently in 5 minutes | You only catch fabrication if you're looking |

**Mindset shift:** at strategy scale, you are not consuming the AI's output — you are *auditing* it. The unit of work is the prompt + the validation pass.

## Four Failure Modes (the ones that bite)

| Failure                     | What it looks like                                                | Fix                                                           |
|-----------------------------|-------------------------------------------------------------------|---------------------------------------------------------------|
| <b>Theme inflation</b>      | 3 distinct complaints lumped under "usability"                    | Demand verbs, not categories; cap themes at 5–7               |
| <b>Citation invention</b>   | "Ticket #4" doesn't exist                                         | Verify every cited source exists; demote uncited claims       |
| <b>Premature confidence</b> | Plausible synthesis with no acknowledgment of gaps                | Force a "what's missing" prompt as second turn                |
| <b>Source bias</b>          | The AI weights the loudest source highest, regardless of accuracy | Ask: "How would your synthesis change if I removed Source X?" |



# The Validated Synthesis Prompt

Role: Research analyst supporting a TPM.

Context: Below are [N] interview transcripts and [M] support tickets from [user segment]. I am investigating [problem statement].

Task: Extract the top 5 recurring pain themes.

Constraints:

- Each theme must cite 2+ specific source identifiers (e.g., "Interview 1, Ticket #4")
- Themes must be verbs, not categories
- Flag any theme you inferred without direct evidence
- At the bottom, list 3 questions the sources do NOT answer

Format: Markdown list; each entry: Theme / Evidence / Confidence (H/M/L) /

Inferred? (Y/N). End with "What's missing" section.

# Activity 1 — Multi-source Synthesis Drill

Triads • 35 minutes

Warm-up: a small healthcare-scheduling pack (different domain). Single-prompt attempt, then iterate after observing failure modes.

- 10 min: single-prompt attempt + run
- 10 min: inspect — which of the 4 failure modes appeared?
- 15 min: iterate the prompt; re-run; compare



10 attempt · 10 inspect · 15 iterate



# Block 2 — The FieldPulse Research Pack

Real materials, real noise, real strategy implications

## What's in the Pack

- 3 anonymized dispatcher interview transcripts (~5 pages each)
- 14 support tickets (already familiar from Week 1 Day 4)
- 1 industry analyst report (5 pages)
- 2 competitor product tour transcripts
- 1 internal "voice of customer" Slack export (50 messages)

**Be honest:** the pack is noisy, partly contradictory, and has gaps. That's the point. Real strategy work feels like this.



# The Cross-Validation Step

Each triad member takes a different theme produced by the AI and verifies **every** citation against the source. Mark any that fail.

- If a citation can't be verified → **cut the citation**, lower theme confidence to L
- If 2+ citations cut from one theme → **cut the theme** or relabel as "hypothesis to validate"
- If a quote is paraphrased misleadingly → **flag** in provenance log

**Don't trust:** the AI's confidence rating about its own output. Independent verification is the only signal.

## Activity 2 — The Research Pack Pass

Triads • 40 minutes

Use your iterated prompt on the real FieldPulse Research Pack. Cross-validate every citation.

- 5 min pre-flight (confirm prompt has constraints)
- 15 min run (split pack across 2–3 calls if needed)
- 15 min cross-validate (one theme per member)
- 5 min edit (cut failed themes; lower confidence)



5 prep • 15 run • 15 verify • 5 edit

# Block 3 — The Competitive Snapshot

The brightest line in AI-assisted research

# ☒ The Bright Line

**AI can summarize and structure source material you provide.  
AI cannot tell you what your competitors actually do.**

## ✓ Safe (you provide source)

- "Structure these two product tour transcripts into a comparison matrix"
- "Identify the 3 most important questions these materials don't answer"
- "Tell me what's missing relative to my own product"

## ✗ Unsafe (AI invents)

- "What is ServiceTitan's market share?"
- "Who is the leader in dispatch SaaS?"
- "Cite a study showing dispatcher productivity gains"



# The Two-Prompt Pattern

## Prompt 1 — Structure what we know:

```
Role: Competitive analyst.  
Context: Pasting in 2 competitor product tour transcripts plus  
their public pricing pages.  
Task: Build a comparison matrix.  
Constraints:  
- Use only facts from the materials I pasted  
- Do NOT pull from your training data  
- Flag any cell where the source is silent  
Format: Markdown table, with citation column.
```

## Prompt 2 — Tell us what's missing:

```
From the matrix above, what 3 questions about these competitors  
do the source materials NOT answer? For each, suggest where a  
TPM should look (analyst reports, customer interviews, public  
earnings calls).
```

## Activity 3 — Competitive Snapshot

Triads • 40 minutes

Use the two-prompt pattern on the competitor materials in the Research Pack. Build a matrix and a list of unanswered questions.

- 15 min: structure-what-we-know prompt + verify cells
- 15 min: what's-missing prompt + decide which gaps matter
- 10 min: write the Competitive Snapshot section of your brief



15 structure • 15 missing • 10 write



# Block 4 — Brief Assembly + Provenance

Stitch into a brief, write the parts AI cannot, log the parts AI did

# The Strategy Brief Structure

1. **Executive summary** — hand-written, 5 bullets max
2. **Top 5 user pain themes** — AI-summarized, evidence-tagged, cross-validated
3. **Competitive snapshot** — AI-drafted from your materials, hand-edited
4. **Strategy implications** — hand-written; ties themes to NS and design principles
5. **Provenance log** — mandatory; documents every AI use and validation step

**Hand-written sections matter most:** they are where the TPM's judgment shows. AI does the volume; you do the meaning.



# The Provenance Log

A trust-building artifact: documents *where* AI was used, *which prompts*, *what validation* was done, and *where the AI was wrong*.

```
## Provenance log
```

```
<h3 id="ai-generated-sections">AI-generated sections</h3>
```

```
<table>
```

```
<thead>
```

```
<tr>
```

```
<th>Section</th>
```

```
<th>Prompts (link)</th>
```

```
<th>Validation</th>
```

```
<th>Findings</th>
```

```
</tr>
```

```
</thead>
```

```
<tbody><tr>
```

```
<td>Pain themes</td>
```

```
<td>Pattern Library #3</td>
```

```
<td>All citations verified</td>
```

```
<td>Cut 1 hallucinated theme</td>
```

```
</tr>
```

```
<tr>
```

```
<td>Competitive matrix</td>
```

```
<td>Pattern Library #5</td>
```

```
<td>Manual check vs source</td>
```

```
<td>Flagged 2 silent cells</td>
```

```
</tr>
```

```
<tr>
```

```
<td>Missing questions</td>
```

```
<td>Pattern Library #6</td>
```

```
<td>Triad consensus</td>
```

```
<td>Adopted 2 of 3</td>
</tr>
</tbody></table>
<h3 id="where-ai-was-wrong">Where AI was wrong</h3>
<ul>
<li>lumped "reconcile" and "timecard" into single theme; we split them</li>
```

```
<li>Executive summary</li>
<li>Strategy implications</li>
```

- This log

# Activity 4 — Brief Assembly + Provenance

Triads • 45 minutes • Readouts

Stitch the day's outputs into a single brief. Hand-write the executive summary and strategy implications. Complete the provenance log.

- 15 min: hand-write executive summary
- 10 min: hand-write strategy implications
- 10 min: complete provenance log
- 10 min: polish for sharing



15 exec · 10 strategy · 10 log · 10 polish · 90s readouts



# Day 4 Wrap

## What we built

- Validated synthesis prompt (Pattern Library entry)
- Two-prompt competitive pattern
- Cross-validation discipline
- AI-Assembled Strategy Brief with provenance

## What opens tomorrow

- **Mapping customer journeys to features**
- Friday readouts: full Week-2 strategy package
- Bridge to Week 3 PRDs

**Bring to Day 5:** Your Strategy Brief, NS, and 3 design principles. All three feed the journey map directly.

# Questions & Reflection

Where in your brief did the AI's confidence exceed your evidence?