Day 2 — Writing Granular Acceptance Criteria

> **Activity packet** for participant triads. Today's job: take your Sections 1–5 draft from Day 1 and write 8–12 testable Acceptance Criteria in Given/When/Then form. By 16:00, the AC section is shareable; tomorrow we add NFRs.

Where we are in the week

Day 1 produced sections 1–5 of the PRD. Today produces section 6 — the Acceptance Criteria. AC are where most PRDs go from "interesting idea" to "implementable contract." A good AC section is the difference between an engineer scoping confidently and an engineer raising 30 questions in standup.

The non-AI rule (still in force)

No generative tools today. Spell-check is fine. Last-week artifacts and Day-1 PRD draft are inputs.

What an Acceptance Criterion is (and isn't)

An AC is a **testable assertion** about the system's behavior in a specific situation. It has three parts:

```
Given <a precondition / starting state>
When <an action or event happens>
Then <an observable, falsifiable result>
```

Is	Is not
Testable (you can write a test that passes/fails)	"The system should be intuitive"
Specific to one behavior	"The system handles errors well"
Implementation-agnostic	"Use Redis to cache the response"
Falsifiable (you can prove it wrong)	"Performance should be good"

The 5 AC failure modes (today's mental model)

Failure	What it looks like	Fix
1. Vague ("intuitive", "fast", "easy")	"Then the user has a good experience"	Replace with measurable outcome
2. Untestable (no observable result)	"Then the user feels confident"	Anchor to a system-visible action or state
3. Restating the goal	"Then dispatchers reconcile faster"	Pin to a specific in-system event, not a metric
4. Multi-condition (AND-soup)	"Then X and Y and Z and W happen"	Split into multiple ACs
5. Implementation-prescriptive	"Then a Redis cache hit returns…"	Describe behavior, not implementation

A good triad will produce AC that fail **none** of these five failure modes.

The "happy path / sad path / weird path" coverage

A complete AC section covers three categories of scenario:

Path	Question	Typical AC count
Happy	What does success look like end-to-end?	3–5
Sad	What happens when the user does the wrong thing or input is invalid?	2–4
Weird	What about network drops, partial data, race conditions, edge cases the user wouldn't think of?	2–4

Triads should aim for **8–12 AC total** with this coverage shape.

The "weird path" is where TPMs earn their seat. PMs who don't think technically miss it. Engineers who don't think about users miss it differently.

Activity 1 — AC Triage

Format: Triad • **35 min** • Block 1

Purpose

Calibrate the eye for the 5 failure modes by triaging real-world AC examples before drafting your own.

Setup

Each triad receives the **AC Triage Pack**: 12 AC examples drawn from real (anonymized) PRDs. Some are good. Some have one of the 5 failures. A few have multiple.

Triad protocol

1. For each AC: identify which failure mode(s) are present (or "clean")
2. For each failed AC: rewrite it
3. Pick the most common failure mode in your pack — be ready to explain why it's the most common in real PRDs

Sample items in the triage pack

- `Given the user is logged in, When they reach the dashboard, Then it loads quickly.`
- `Given a dispatcher with 3+ open tickets, When they tap "Reconcile all", Then the modal opens with all 3 tickets pre-selected and shows the count "3 selected" in the header.`
- `Given any user, When they use the system, Then they should feel productive.`
- `Given a network drop, When the user submits the form, Then the data is queued via WebSocket retry mechanism.`
- `Given a duplicate ticket, When the API receives it, Then 409 Conflict is returned with the original ticket ID in the body.`

Readout (60 seconds per triad)

> "The most common failure in our pack was [X]. Our cleanest rewrite was [example]."

Deliverable

12 triaged AC with failure-mode labels, plus rewrites for each failing AC and a one-sentence note on the most common failure.

Activity 2 — Happy-Path AC for Your PRD

Format: Triad • **40 min** • Block 2

Purpose

Write the 3–5 happy-path AC for the triad's PRD. These are usually the easiest — get them out of the way first.

Setup

Each triad needs their Section 5 solution sketch from Day 1, the AC template card, and the 5-failure-mode checklist. No AI.

Triad protocol

1. **Identify the happy path** (5 min). Re-read Section 5 from yesterday. The 4–8 step flow becomes the basis for AC.
2. **Solo drafts** (15 min). Each member writes 3 happy-path AC alone.
3. **Pool, de-dupe, refine** (15 min). The triad converges on 3–5 final happy-path AC.
4. **Failure-mode check** (5 min). Run each AC through the 5-failure list.

What "good" looks like

- Each AC is one Given/When/Then with no ANDs in the Then
- Each AC names the exact system state, screen, or event
- A junior engineer could read the AC and write a test against it

Deliverable

3–5 happy-path AC in Given/When/Then form, each passing the 5-failure-mode check, appended to PRD Section 6.

Worked FieldPulse example (happy path)

```
Given a dispatcher viewing the end-of-shift summary
When they tap "Reconcile all"
Then the reconcile modal opens within 1 second
  and shows all open tickets pre-selected
  and the header shows the count of tickets selected
```

(The "and"s here are part of one observable state, not separate behaviors. This is a judgment call — split if the "ands" represent different system actions.)

Activity 3 — Sad-Path and Weird-Path AC

Format: Triad • **40 min** • Block 3

Purpose

Cover the failure modes — what goes wrong when the user, the network, or the data does something unexpected.

Setup

Each triad needs the happy-path AC from Activity 2 and the sad/weird-path generator prompts. No AI.

The sad-path generator

For each happy-path AC, ask:

- What if the **input is invalid**?
- What if the **user lacks permission**?
- What if the **user cancels mid-flow**?
- What if the **data is missing or stale**?

The weird-path generator

For each happy-path AC, ask:

- What if the **network drops** mid-action?
- What if **two users do the same thing simultaneously** (race)?
- What if the **action takes longer than expected** (timeout)?
- What if the **upstream system is down**?
- What if the **user is at the boundary** (max characters, max items, empty)?

Triad protocol

1. **Solo sad-path drafts** (10 min). Each member produces 2 sad-path AC for the triad's PRD.
2. **Solo weird-path drafts** (10 min). Each member produces 2 weird-path AC.
3. **Pool and cull to 4–6** (10 min). 2–4 sad, 2–4 weird.
4. **Failure-mode check** (10 min). The 5 failure modes again.

Worked FieldPulse examples

Sad path:

```
Given a dispatcher with 0 open tickets
When they tap "Reconcile all"
Then a non-modal toast appears with text "No tickets to reconcile"
  and the dispatcher remains on the summary screen
```

Weird path:

```
Given a dispatcher with 47 open tickets (more than the modal can list)
When they tap "Reconcile all"
Then the modal opens with the first 25 tickets pre-selected
  and a banner reads "22 more tickets — load more"
  and a tap on the banner appends the next 25 below
```

Deliverable

2–4 sad-path AC and 2–4 weird-path AC appended to PRD Section 6, each passing the 5-failure-mode check.

Activity 4 — AC Cross-Review

Format: Triad-pair • **45 min** + Wrap • Block 4

Purpose

Apply the calibrated eye to *another triad's* AC. Cross-review surfaces failure modes the author triad has gone blind to.

Setup

Instructor pre-assigns triad pairs. Each triad needs their own Sections 1–6 (incl. AC) and a printed review template. No AI.

The protocol

1. **Pair triads** (instructor assigns).
2. **Swap AC sections** (5 min). Read the other's PRD Sections 1–5 first for context, then their AC.
3. **Review pass — identify failures** (15 min). For each AC:
    - Which failure mode (if any)?
    - Which path (happy / sad / weird) is *missing*?
    - Is there an implicit AC the author triad didn't write?
4. **Author triad responds** (10 min). For each finding: adopt / defer / push back.
5. **Authors revise** (15 min). Update the AC section.

Output

- A revised AC section (8–12 ACs)
- A short **review-resolution note** at the bottom of Section 6 listing what was adopted, deferred, or rejected (and why)

Wrap (last 15 min)

Each triad shares **one AC** they're proudest of and **one** they're least sure about. The "least sure" share is the calibration — what does the room think about it?

End-of-day checkpoint

Each triad leaves the day with their PRD updated to include:

- [x] Section 6 with **8–12 testable AC** in Given/When/Then form
- [x] Coverage spans happy / sad / weird paths
- [x] All AC pass the 5-failure-mode check
- [x] A review-resolution note documenting changes from peer review