



Week 3 Knowledge Check

Requirements Engineering & Mini-Capstone

Technical Product Manager Academy



How This Works

12 questions across the week — technical PRDs, acceptance criteria, non-functional requirements, PRD assembly, and structured review.

Commit to an answer **before** the reveal. Retrieval is the point.

? Question 1

What is the "TPM contract" — the test that distinguishes a Technical PRD from a generic PM PRD?

- A) A PRD the executive sponsor approves without edits
- B) A PRD an engineering lead would accept as scoping input *without a clarifying call*
- C) A PRD that fits on a single page
- D) A PRD marketing can turn into launch copy



Answer 1 — B

An engineering lead would accept it as scoping input without a clarifying call.

The single best evaluation criterion — and the bar Friday's review uses.

? Question 2 — True / False

Week 3 encourages using AI to draft the PRD faster, so triads can spend more of their time on review.

Answer 2 — False

Week 3 is deliberately **non-AI**.

AI fills vagueness with confident, generic prose — and vagueness is the bug we're eliminating.

? Question 3

Which of these *belongs* in the §5 Solution Sketch, rather than being cut as over-specification?

- A) Database schema decisions
- B) Specific API contracts
- C) The user-visible flow (4–8 steps)
- D) Class names and library picks



Answer 3 — C

The user-visible flow (4–8 steps).

Schema, API contracts, and class names are engineering judgment — leave room for it.

? Question 4

An acceptance criterion reads: *"Then dispatchers reconcile faster."* Which of the five AC failure modes is this?

- A) Vague
- B) Untestable
- C) Restating the goal
- D) AND-soup



Answer 4 — C

Restating the goal.

It pins to a metric, not an in-system event. Fix by anchoring to an observable system state.

? Question 5 — True / False

Weird-path AC coverage — network drops, races, timeouts, boundary values — is described as the TPM differentiator that PMs and pure engineers each tend to miss.

Answer 5 — True

Weird-path coverage is the TPM differentiator.

PMs who don't think technically miss it; engineers who don't think about users miss it differently.

? Question 6

In the Requirement / Defense / Verification NFR form, which field is singled out as the one that separates a real TPM NFR from boilerplate?

- A) Requirement
- B) Defense
- C) Verification
- D) Category



Answer 6 — B

Defense.

"p95 < 400ms because dispatchers tap this 40 times per shift" beats "should be fast."

? Question 7

How should a performance NFR express its latency target?

- A) The average (mean) response time
- B) A p95 target on a representative device and connection
- C) A best-case time under ideal conditions
- D) "Fast enough that users don't complain"



Answer 7 — B

A p95 target on a representative device/connection.

Averages lie: a bimodal 200ms / 5s split averages to 2.6s but describes no real user's experience.

? Question 8 — True / False

A mature NFR section includes a "Known Trade-offs" subsection that names the tensions — and their resolution or open status — rather than pretending everything is achievable at once.

Answer 8 — True

The "Known Trade-offs" subsection is the senior-author signal.

"We accept this" is strong. "We'll change it if X happens" is stronger.

? Question 9

In §10 Dependencies, what makes a dependency "owned"?

- A) It's assigned to the responsible team
- B) It's assigned to a named person
- C) Its status is marked "Confirmed"
- D) It has a "by when" date



Answer 9 — B

A named person.

"The auth team" is unowned; "Lin from auth" is owned. "TBD" as owner is the dependency that ships late.

? Question 10 — True / False

In §9, listing "no risks" is acceptable as long as the feature is genuinely low-risk.

Answer 10 — False

"No risks" is a hard fail.

It's the tell of a PRD written to avoid scrutiny. Each risk gets likelihood, impact, and a mitigation.

? Question 11

On the Friday review rubric, which dimension carries the highest weight?

- A) Problem clarity (20%)
- B) AC testability (25%)
- C) NFR completeness (20%)
- D) Strategy linkage (15%)

Answer 11 — B

AC testability — 25%, the highest weight.

Score ≥ 3.0 ships as-is; 2.0–2.9 ships with named gaps; < 2.0 triggers facilitator intervention.

? Question 12

Which reviewer rule is enforced during AC cross-review and again on Friday?

- A) Always give three specific strengths before any criticism
- B) If you say "this could be more specific," you also write the specific version
- C) Never score below 2.0 without facilitator sign-off
- D) Reviewers must not coordinate with each other

Answer 12 — B

"If you say it could be more specific, you write the specific version."

The rule that eliminates lazy review feedback. Enforce it.



That's Week 3

Score yourself — but the real check is whether your PRD would survive a hostile engineer on Monday.