

SLO Triage Pack — 8 Statements

Day 4 · Activity 1 handout. Eight SLO statements drawn from real (anonymized) systems. Triage each: is it clean, or does it fail one of the common checks? Then rewrite the five worst so they pass **percentile + window + defense**.

The common failure modes

Failure mode	What it looks like
No measurement window	A target with no period — "always"? "best case"?
No percentile (or "average")	Averages hide bimodal distributions
No defense	A number with no rationale for <i>why that number</i>
Aspirational beyond capability	A target the team's staffing/infra can't sustain
Mismatched to user threshold	A target far looser or tighter than the user actually needs
Confusion with SLA	A <i>contractual</i> commitment stated as an internal target

The 8 statements

1. `p95 < 400ms`
2. `Average response time < 500ms`
3. `99.9% availability`
4. `99.99% availability` (for a 2-engineer feature on shared infra)
5. `p95 < 5s` (for a typing-feedback feature)
6. `We commit to 99.99% availability`
7. `p99 ≤ 800ms, 99.9% of the time over 30 days`
8. `Error rate under 1% during business hours, measured on the checkout endpoint`

Your protocol

1. **Triage all 8** (15 min). Label each with its failure mode(s), or mark "**clean**" if it passes.

2. **Rewrite the 5 worst** (15 min) so each passes all three checks: **percentile + window + defense**.
3. **Identify the one with the most subtle failure** (5 min) — discuss in the readout.

The three checks a clean SLO passes

- **Percentile** — p95 / p99, *not* "average."
- **Window** — "over 30 days," *not* "always" / unstated.
- **Defense** — a stated reason tied to user behavior or platform precedent, *not* a bare number.

Readout (60 sec)

"The most subtle failure was [X] because [why]. Our cleanest rewrite was [example]."

Triage grid

#	Statement	Failure mode(s) / "clean"	Rewrite (if among worst 5)
1			
2			
3			
4			
5			
6			
7			
8			