



Monolith vs Microservices

Week 4 • Day 1 • Technical Architecture & Constraints

TPM Academy



Day 1 Objectives

- Frame mono-vs-micro decisions in **business + operational terms**, not framework hype
- Apply the **three-question frame**: deploy cadence, failure domain, scaling axis
- Recognize the "**majestic monolith**" as the often-correct default
- Draft **Section 1 of the TCD** — your architecture stance with trade-off and revisit trigger
- Begin **Section 2 — Integration map** with failure-handling stances



Today's Shape

Clock	Block	Outcome
09:00 – 09:15	Opening	AI restored; PRD on the wall
09:15 – 10:30	Block 1 + Activity 1	Mono/Micro Triage
10:45 – 12:00	Block 2 + Activity 2	Apply the frame to your PRD
13:00 – 14:15	Block 3 + Activity 3	Integration map (first pass)
14:30 – 16:00	Block 4 + Activity 4 + Wrap	AI-assisted architecture Q&A

Block 1 — The Vocabulary & the Frame

Enough to follow the architect's reasoning, not enough to win arguments by quoting taxonomy



The Vocabulary Card

Term	What it means	TPM signal
Monolith	One deployable, one repo, one runtime	Faster start; harder team scaling
Modular monolith	Monolith with strict module boundaries	Often the right answer
Microservice	Independently deployable, owned by one team	Independent scale & failure domains
SOA	Coarse-grained services, often shared infra	What "microservices" often actually is
Distributed monolith	Multi-deploy that <i>must</i> deploy together	Anti-pattern
FaaS	Per-function deploys	Bursty workloads; not a default
Event-driven	Components via queues / events	Decouples timing; complicates debugging

? The Three-Question Frame

1. Deploy cadence

Whose deployment cadence does separation protect?

If only one team owns it, separation buys mostly ops cost.

2. Failure domain

What independent failure do we want?

If the feature can't function without its dependencies anyway, separation buys no resilience.

3. Scaling axis

What scale curve are we anticipating?

If load is uniform, one runtime is cheaper.

Heuristic: 2+ "yes with evidence" → microservice plausible. 2+ "we don't know" → modular monolith.



The "Majestic Monolith" Position

Most early-stage features ship better as **modules in a well-factored monolith** than as new services.

- **Faster end-to-end testing** (one process, one deploy)
- **Lower observability complexity** (no inter-service tracing tax)
- **Lower coordination cost** (no shared API contract)
- **Reversibility** — extract a module to a service later; merging two services is hard

The TPM job: surface the business cost of *premature* service-orientation. Not to win the debate — to set the right starting point.

Activity 1 — Mono/Micro Triage

Triads • 35 minutes

8 anonymized examples. Apply the three-question frame; stamp Monolith / Micro / Modular monolith / Need more info.

- 15 min: stamp all 8
- 10 min: pick 2 strongest cases for separation, 2 for staying together
- 10 min: identify 1 disagreement to share



15 stamp · 10 pick · 10 disagree

Block 2 — Apply the Frame to Your PRD

Draft TCD §1 — Architecture stance



The Stance Template

```
## 1. Architecture stance
```

```
<p>We propose this feature ship as <strong><modular monolith / new module  
in the existing service / new microservice / hybrid></strong> because:</p>  
<ul>  
<li><strong>Deploy cadence:</strong> [whose tempo are we protecting]</li>  
<li><strong>Failure domain:</strong> [what should keep running if X fails]</li>  
<li><strong>Scaling axis:</strong> [what curve do we expect]</li>  
</ul>
```

```
<strong>Trade-off accepted:</strong> [the cost of this choice – be honest]  
<strong>Revisit if:</strong> [the trigger that would change the answer]
```

"Revisit if" is the senior-author signal. A concrete trigger ("if our deploy frequency drops below daily" or "if traffic to this feature grows 10x") beats "later."



Worked Stance (FieldPulse reconcile)

1. Architecture stance

<p>We propose this feature ship as a new module in the existing mobile-app monolith because:</p>

Deploy cadence: Only the Dispatcher Workflow squad will own this code; we have no cross-team deploy contention to resolve.

Failure domain: Reconcile depends on Tickets API and Auth; if either is down, the feature can't function. Separation buys no resilience.

Scaling axis: Reconcile traffic is bounded by dispatcher shifts (predictable daily curve, < 10K active dispatchers). Same scale curve as the rest of the mobile app.

<p>Trade-off accepted: A future split (e.g., into a "shift-end workflows" service) becomes harder once this code lives next to existing modules.</p>

Revisit if: Reconcile traffic exceeds 5x rest-of-app, OR a second team starts contributing reconcile features.

Activity 2 — Apply the Frame

Triads • 40 minutes

Run the three-question frame against your locked PRD feature. Draft TCD §1.

- 5 min: re-read PRD §10 (Dependencies)
- 15 min: 3-question pass; honest answers with evidence
- 15 min: draft §1
- 5 min: sanity check — real evidence or hand-wave?



5 + 15 + 15 + 5

Block 3 — Integration Map (First Pass)

What the feature touches; how it fails



The Integration Table

System	Owner	Sync/Async	R/W	Failure handling
Auth (SSO)	Identity	Sync	R	Fail closed
Audit event store	Platform	Async	W	Queue + DLQ
Tickets API	Dispatch	Sync	RW	Retry x3, then fail
Notification svc	Comms	Async	W	Best-effort, drop

"Failure handling: TBD" is fine if it's an open question. **"Hopefully it works"** is not.

Two-Way Contracts

Some integrations require **coordination with another team** — they have to change something for you. Mark these explicitly.

Contract	What we need	What they need
Tickets API	New endpoint for batch close	Schema review; rate-limit budget
SSO scopes	New scope `reconcile.write`	Security review; rollout plan
Audit schema	7 new event types	Schema versioning; PII review

Two-way contracts go on the **Stakeholder + Sign-off matrix** (TCD §6, Friday).

Activity 3 — Integration Map

Triads • 40 minutes

List integrations from PRD §10 + anything missed. Characterize each. Mark two-way contracts.

- 15 min: list integrations
- 15 min: characterize sync/async, R/W, failure
- 10 min: identify two-way contracts



15 + 15 + 10

Block 4 — AI is Back

Stress-test the stance; surface what we forgot



AI Norms (Refresher)

- Use Week 1 prompt patterns: **Role / Context / Constraints / Format**
- Use Week 2 Day 4 discipline: **provenance log** for any AI-assisted output
- Stay on the safe side of the bright line: **structure source material; do not ask AI for facts about your real systems**
- Cross-validate: AI surfaces possibilities; you decide what's real

Not allowed: "AI told me our system uses X." That's hallucination. AI doesn't know your codebase.



Today's Two Prompts

A. Stress-test our stance

Paste TCD §1 + integration table. Ask for the 3 strongest objections, each with the specific scenario.

B. What did we forget?

Paste integration table. Ask: "what categories of dependency are commonly forgotten for a feature with this footprint?"

```
Role: Senior staff engineer reviewing an architecture stance.  
Context: <paste tcd="§1" += "integration=" table=" ">  
Task: Identify the 3 strongest objections to the stance.  
Constraints:  
- Treat the proposed approach as a starting point, not a verdict  
- For each objection, name the specific scenario where it would matter  
- Flag any place the rationale relies on assumptions not stated above  
Format: Numbered list of objections, each with a scenario.</paste>
```

Activity 4 — AI-Assisted Architecture Q&A

Triads • 45 minutes • Wrap

Run Prompt A and Prompt B. Adopt / defer / reject each suggestion. Update §1 + §2. Add provenance note.

- 10 min: Prompt A → 3 objections
- 10 min: adopt / defer / reject decisions
- 10 min: Prompt B → missing dependencies
- 10 min: update §§1–2 + provenance note
- 5 min: AI-prose check; rewrite if generic



10 + 10 + 10 + 10 + 5 · 15-min wrap



Day 1 Wrap

What we built

- Three-question frame (deploy / failure / scale)
- Architecture stance with revisit trigger
- Integration table with failure handling
- Two-way contracts list
- AI provenance discipline restored

What opens tomorrow

- **System security & compliance**
- STRIDE threat-model pass
- Revising PRD's Security NFRs at architectural level

Bring to Day 2: Your TCD §§1–2. We'll thread the security model through them tomorrow.

Questions & Reflection

If your stance is wrong, which assumption fails first?