

Storage-Choice Buckets

Day 1 handout. The four storage families most B2B features choose between — and when to reach for each. Your job is not to pick the engine; it's to describe the access pattern clearly enough that the architect can.

Bucket	When to reach for it
Relational (Postgres, MySQL)	Strong relationships, transactions, complex queries; default for most B2B
Document (MongoDB, Firestore)	Flexible schema, hierarchical data; good when shape varies per record
Key-value (Redis, DynamoDB)	High throughput, simple lookups, caches, session state
Search (Elasticsearch, OpenSearch)	Full-text search, analytics on text

How to use this on your feature:

1. Describe the dominant access pattern (from your Access Pattern Sheet).
2. Match it to the bucket that serves it cheapest.
3. Note where a second bucket might carry a secondary pattern (e.g., relational primary + search index for full-text).

A TPM **describes the access pattern**; the architect **picks the engine**. Don't skip straight to the engine name.