



Database Structures & Data Logic

Week 5 • Day 1 • Technical Infrastructure & Modeling

TPM Academy



Day 1 Objectives

- Read a data model at the entity level: **keys, indexes, relationships, normalization vs denormalization**
- Apply the "**access pattern first**" discipline — queries before schema
- Make the **storage choice** argument in customer terms
- Surface **schema-level trade-offs** using the Week-4 template
- Draft **Technical Modeling Document (TMD) Section 1 — Data model**



Today's Shape

Clock	Block	Outcome
09:00 – 09:15	Opening	Pin TCD Sections 1–6 on the wall
09:15 – 10:30	Block 1 + Activity 1	Read & write access patterns
10:45 – 12:00	Block 2 + Activity 2	Draft the entity model
13:00 – 14:15	Block 3 + Activity 3	Storage trade-offs
14:30 – 16:00	Block 4 + Activity 4 + Wrap	AI critique + polish

Block 1 — Access Patterns First

Queries before schema



The Vocabulary Card

Term	What it means
Entity	A "thing" the system tracks (Dispatcher, Ticket, ReconcileEvent)
Primary key	The unique identifier for an entity
Foreign key	A reference from one entity to another
Index	Speed up reads at the cost of writes + storage
Normalization	Splitting data to avoid duplication
Denormalization	Duplicating data on purpose for read speed
Cardinality	1:1, 1:N, N:N — the relationship shape
Constraint	A database-enforced rule (NOT NULL, UNIQUE, CHECK)



The Four Storage Buckets

Bucket	When to reach for it
Relational (Postgres, MySQL)	Strong relationships, transactions, complex queries — default for B2B
Document (MongoDB, Firestore)	Flexible schema, hierarchical data; shape varies per record
Key-value (Redis, DynamoDB)	High throughput, simple lookups, caches, sessions
Search (Elasticsearch, OpenSearch)	Full-text search, analytics on text

The TPM doesn't pick the engine. The TPM describes the access pattern clearly enough that the architect can.

The "Access Pattern First" Discipline

1. **List the queries the feature needs.** "Show me a dispatcher's open tickets, sorted by latest activity."
2. **For each query: what does it read, in what order, with what filter?**
3. **Now design the schema** so the queries are cheap.

If your data model can't be defended by the queries that drive it, it's wrong.

Activity 1 — Read & Write Patterns

Triads • 35 minutes

List every read and write the feature needs. Capture data, frequency, freshness requirement.

- 5 min: re-read Product Requirements Document (PRD) Section 5
- 15 min: list reads (data, filter, freshness)
- 10 min: list writes (entity, invariants, concurrency)
- 5 min: highlight 1–2 highest-volume reads



5 + 15 + 10 + 5

Block 2 — Draft the Entity Model

Tables, keys, indexes, relationships



The Entity Template

```
## Entity: <Name>

| Field | Type | Notes |
|-----|-----|-----|
| id     | UUID | Primary key |
| ...    | ...  | NOT NULL? UNIQUE? FK to ...? |

**Indexes:**
- (field, field) – purpose: query #N from access patterns

**Relationships:**
- belongs_to <Entity> via <fk_field>
- has_many <Entity> via <fk_field>

**Notes / invariants:**
- A reconcile_event is immutable once created.
- ticket.status ∈ {open, reconciled, voided}.
```

Indexes reference access patterns by number. "Index for query #4" beats "useful index."



Worked Entity (FieldPulse)

```
## Entity: ReconcileEvent

| Field | Type | Notes |
|-----|-----|-----|
| id | UUID | PK |
| dispatcher_id | UUID | FK → Dispatcher |
| ticket_ids | UUID[] | The set reconciled (denormalized) |
| submitted_at | timestamp | NOT NULL |
| client_ip | inet | for audit |
| user_agent | text | for audit |
| signature_hash | text | for non-repudiation (R-threat) |

**Indexes:**
- (dispatcher_id, submitted_at DESC) – serves pattern #4 (history)
- (submitted_at) – serves pattern #3 (audit-trail range query)

**Relationships:**
- belongs_to Dispatcher via dispatcher_id

**Invariants:**
- Immutable: never UPDATE'd after INSERT
- ticket_ids is denormalized (see trade-off #2)
```

The `signature_hash` addresses the Repudiation threat in the STRIDE model (Spoofing, Tampering, Repudiation, Information disclosure, Denial of service, Elevation of privilege).

Activity 2 — Entity Model

Triads • 40 minutes

List entities. For each, draft fields, indexes, relationships, invariants.

- 10 min: list entities (new + existing)
- 15 min: draft fields per new entity
- 10 min: indexes (referencing patterns)
- 5 min: relationships + cardinalities



10 + 15 + 10 + 5

Block 3 — Storage Trade-Offs

Use the Week-4 trade-off template at schema scale

Three Schema-Level Trade-Offs

Trade-off	Tension
Normalization vs denormalization	Clean writes vs fast reads
Strong consistency vs replica reads	Always-fresh vs scalable read load
Schema flexibility vs schema integrity	Easy migration vs caught-by-database integrity

Same Week-4 template: Option A vs Option B / Choice / Why / Accepted cost / Revisit trigger.



Worked Schema Trade-Off

```
### Trade-off – ticket_ids storage
**Tension:** Normalization vs denormalization.
**Option A:** ReconcileEvent.ticket_ids: UUID[] (denormalized).
**Option B:** Separate ReconcileEventTicket join table.
**Choice:** Option A.
**Why:** Reconcile events are read in batches of 1 most of the
time. Reading the array is one row. Option B requires
a join. Reconcile events are immutable and rarely
queried by ticket – access patterns don't reward
normalization.
**Accepted cost:** Cannot index by individual ticket_id (no
"find every reconcile that touched ticket X" query
without a denormalization).
**Revisit trigger:** If "find reconciles touching ticket X"
becomes a high-volume query.
```

Activity 3 — Storage Trade-Offs

Triads • 40 minutes


- 15 min: identify 3 trade-off points in your model
- 15 min: write each (Option A/B/Choice/Why/Cost/Trigger)
- 10 min: cross-check against Technical Concept Document (TCD) Section 4 Service Level Objectives (SLOs)



15 + 15 + 10

Block 4 — AI Critique + Polish

Treat the AI as a critic, not an oracle



Today's Two Prompts

A. Critique this schema

Role: Senior backend engineer reviewing a feature's data model.

Context: <entity definitions="" += "" indexes="" relationships="" access="" pattern="" sheet="">

Task: Identify the top 3 issues.

Constraints:

- Name the specific access pattern or invariant that breaks
- Suggest a concrete fix (new index / entity / cardinality)
- Do not suggest generic best practices

Format: Numbered – Issue / What breaks / Fix.

</entity>

B. What did we forget?

Continuing from above. Given these access patterns, what queries the user will eventually want are NOT yet supported by this schema? List 3, ranked by likelihood within 6 months.

Activity 4 — AI Critique

Triads • 45 minutes • Wrap

- 10 min: Prompt A → 3 issues
- 10 min: adopt / defer / reject
- 10 min: Prompt B → forgotten queries
- 15 min: update Section 1 + provenance + AI-prose check



10 + 10 + 10 + 15 · 15-min wrap



Day 1 Wrap

What we built

- Access Pattern Sheet
- Entity model (PKs, indexes, relationships)
- 3+ explicit storage trade-offs
- TMD Section 1 drafted

What opens tomorrow

- **Cloud architecture & infrastructure**
- Region / AZ / managed services
- Multi-tenancy stance

Bring to Day 2: Your TMD Section 1 + the TCD's Container diagram. Tomorrow we put the boxes in regions.

Questions & Reflection

Which of your indexes would you cut if writes became 10x more expensive?