



Week 8 Knowledge Check

AI Spec Development & Capstone

Technical Product Manager Academy



How This Works

12 questions across the week — AI Spec pattern, compressed discovery + PRD, compressed architecture, validation + finalization, and Friday's presentations.

Commit to an answer **before** the reveal. Retrieval is the point.

? Question 1

In the AI Spec Development pattern, how is the 5-prompt sequence run?

- A) All 5 prompts in parallel, then validate once at the end
- B) Each prompt validated before the next runs — a validation chain
- C) As one mega-prompt with every input at once
- D) Prompts 1–4 draft; prompt 5 validates all four together



Answer 1 — B

Each prompt validates before the next runs.

The validation chain is what keeps the AI Spec honest. Running all five in parallel — or as one mega-prompt — breaks it.

? Question 2 — True / False

Senior TPMs use AI Spec Development to draft the spec, not to do the thinking — judgment shapes every section and the AI accelerates the typing.

Answer 2 — True

"Use this to draft, not to think."

Judgment shapes every section; AI accelerates the typing. AI Spec without validation is fiction.

? Question 3

In the capstone, what does "compressed" work mean?

- A) Lower-quality work produced faster
- B) Selective — keep what's load-bearing, cut what's cosmetic
- C) Only the happy paths, skipping edge cases
- D) AI-generated drafts with lighter validation

Answer 3 — B

Selective: keep what's load-bearing, cut what's cosmetic.

Compressed work is *not* lower-quality work. Every word earns its place.

? Question 4

In the compressed PRD-light, what's the acceptance-criteria coverage target?

- A) 6 happy / 0 sad / 0 weird
- B) 3 happy / 2 sad / 1 weird
- C) 2 happy / 2 sad / 2 weird
- D) 5 happy / 1 sad / 0 weird



Answer 4 — B

3 happy / 2 sad / 1 weird.

Compressed means *fewer* ACs, not weaker ones — still Given/When/Then, same 5-failure-mode discipline.

? Question 5

Running the AI Spec sequence under time pressure, which content must be preserved *verbatim* rather than paraphrased by the AI?

- A) The engineering-ready summary paragraphs
- B) Trade-off costs, revisit triggers, named invariants, customer quotes
- C) The headline and section titles
- D) The provenance log entries

Answer 5 — B

Trade-off costs, revisit triggers, named invariants, and customer quotes.

Everything else can be synthesized — but hallucinated claims get cut. Time pressure tempts skipping validation; hold the line.

? Question 6 — True / False

In the Day-3 cross-triad spot-check, a reviewer triad that reads the AI Spec and reports "looks fine" without checking claims against the source has done its job correctly.

Answer 6 — False

Watch for **rubber-stamping**.

The reviewer should find *at least one* weak claim — checked against the source. "Looks fine" without checking is the failure.

? Question 7

The Day-4 7-point technical-logic validation exists mainly to check what?

- A) That each artifact is well-written on its own
- B) Whether the artifacts fit together as one coherent set
- C) That the AI Spec matches the deck word-for-word
- D) That grammar and formatting are consistent



Answer 7 — B

Whether the artifacts fit together.

Each artifact may be fine alone — the 7 checks ask if they cohere (e.g., the SLO budget fits the sequence diagram; the data model supports the API).

? Question 8 — True / False

Most triads pass all 7 technical-logic checks on the first pass.

Answer 8 — False

Most triads **fail 2–3 checks** on the first pass.

That's expected — that's exactly why the finalization day exists.

? Question 9

In the cross-artifact integration pass, a good cross-reference must be:

- A) Short and informal — "see the TCD" is enough
- B) Specific (cite the §), verified, and add value
- C) Placed only in the AI Spec, not the other documents
- D) Generated by AI to guarantee consistency



Answer 9 — B

Specific, verified, and adds value.

"See the TCD" without a section number is a sign the writer never actually verified it. Cite the §.

? Question 10

On the Friday presentation rubric, which dimension carries the *single highest* weight?

- A) Customer + problem clarity (15%)
- B) AI Spec quality (20%)
- C) Trade-off honesty (15%)
- D) Presentation craft (15%)

Answer 10 — B

AI Spec quality — 20%.

The highest-weighted dimension. The week's signature artifact carries the most weight; the rest cluster at 10–15%.

? Question 11 — True / False

Triads are scored on a curve — ranked against one another so the scores spread out.

Answer 11 — False

Score against the **rubric**, not against other triads.

Whole-number scores, no curve: if two triads earn 3.5, both are 3.5.

? Question 12

During Q&A, when asked a specific question you *cannot* answer, the disciplined response is:

- A) Give a confident best guess so you don't look unprepared
- B) "Honestly, I don't know — here's what we'd do to find out"
- C) Immediately defer the question to the instructor
- D) Restate your customer + problem to fill the time

Answer 12 — B

"Honestly, I don't know — here's what we'd do to find out."

Curiosity, not defensiveness. A well-handled "I don't know" beats a hand-wave — and triads who handle Q&A this way score high on presentation craft.



That's Week 8 — and the Academy

Score yourself — but the real check is whether you can *use* these in your next role. Take the muscles. Keep the discipline.